

Bioinformatics

Lab2

(실습 이해했는지에 대한 문제 하나 내음
근데 쓰느느간 아님)

Index

1. 개체 명 인식
2. 데이터 세트
3. Environment
4. Transformers
5. How to train NER using transformers


개체 명 인식

NER

Language model 정보 추출의 한 프로세스

- 개체 명 인식 (named entity recognition): 미리 정의해 둔 사람, 회사, 장소, 시간, 단위 등에 해당하는 단어(개체 명)을 문서에서 인식하여 추출 분류하는 기법
- 일반적으로 IOB tag를 사용; B-Begin, I-Inside, O-Others

text에서
원하는 entity type을 찾아내는 방법



데이터 세트

- https://huggingface.co/datasets/ncbi_disease

id (string)	tokens (sequence)	ner_tags (sequence)
"0"	["Identification", "of", "APC2", ",", "a", "homologue", "of", "the", "adenomatous", ...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 2, 2, 2, 0, 0]
"1"	["The", "adenomatous", "polyposis", "coli", "(", "APC", ")", "tumour", "-", "suppressor" ...	[0, 1, 2, 2, 2, 2, 2, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
"2"	["Complex", "formation", "induces", "the", "rapid", "degradation", "of", "betacatenin", ...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
"3"	["In", "colon", "carcinoma", "cells", ",", "loss", "of", "APC", "leads", "to", "the", ...	[0, 1, 2, 0, ...
"4"	["Here", ",", "we", "report", "the", "identification", "and", "genomic", ...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
"5"	["Mammalian", "APC2", ",", "which", "closely", "resembles", "APC", "in", ...	[0, ...
"6"	["Like", "APC", ",", "APC2", "regulates", "the", "formation", "of", "active", ...	[0, 1, 2, ...

ner_tags (IOB tag)

- 0 – O
- 1 – B-disease
- 2 – I-disease

Environment


- Colab
 - Python 3.7.13
 - CUDA 11.2
 - Numpy 1.21.6
 - Scikit-learn 1.0.2
- Pytorch 1.12.0
- Transformers 4.24.0
- Datasets 2.7.1
- Evaluate 0.3.0 성능 평가 모듈
- seqeval
정답을 맞아서 정의해주는 라이브러리


Transformers

- <https://huggingface.co/docs/transformers/index>

Transformers


State-of-the-art Machine Learning for PyTorch, TensorFlow, and JAX.

 Transformers provides APIs and tools to easily download and train state-of-the-art pretrained models. Using pretrained models can reduce your compute costs, carbon footprint, and save you the time and resources required to train a model from scratch. These models support common tasks in different modalities, such as:

 **Natural Language Processing:** text classification, named entity recognition, question answering, language modeling, summarization, translation, multiple choice, and text generation.

 **Computer Vision:** image classification, object detection, and segmentation.

 **Audio:** automatic speech recognition and audio classification.

 **Multimodal:** table question answering, optical character recognition, information extraction from scanned documents, video classification, and visual question answering.

How to train NER using transformers

Code: <https://github.com/yesol-park/Bioinformatics2022>

- Get dataset from dataset repository on the Hub

```
train_dataset = load_dataset("ncbi_disease", split="train")
test_dataset = load_dataset("ncbi_disease", split="test")

print(train_dataset)
print(test_dataset)
```

```
Dataset({
  features: ['id', 'tokens', 'ner_tags'],
  num_rows: 5433
})
Dataset({
  features: ['id', 'tokens', 'ner_tags'],
  num_rows: 941
})
```

How to train NER using transformers

- Load tokenizer and pretrained model

```
model_name = "bert-base-uncased"

tokenizer = AutoTokenizer.from_pretrained(model_name)

config = AutoConfig.from_pretrained(
    model_name,
    num_labels=num_labels,
    finetuning_task="ner")
model = AutoModelForTokenClassification.from_pretrained(
    model_name,
    config=config,
)
```

Bert는 token 단위로 나뉘는

- Preprocessing the dataset

```
def tokenize_and_align_labels(examples):
    tokenized_inputs = tokenizer(examples["tokens"], padding="max_length", truncation=True,
                                max_length=256, is_split_into_words=True, )

    labels = []
    for i, label in enumerate(examples["ner_tags"]):
        word_ids = tokenized_inputs.word_ids(batch_index=i)
        previous_word_idx = None
        label_ids = [
            (None, 0, 1, 2, 2, 2, 3, 4, 5, 5, 6, 7, 8, 8, 8, 8, 9, 9, 9, 10, 11, 11, 12, 12, 13, None, None, None, None, None, None, None, None),
            for word_idx in word_ids:
                # Special tokens have a word id that is None. We set the label to -100
                # so they are automatically ignored in the loss function.
                if word_idx is None:
                    label_ids.append(-100)
                # We set the label for the first token of each word.
                elif word_idx != previous_word_idx:
                    label_ids.append(label_to_id[label[word_idx]])
                # For the other tokens in a word, we set the label to -100.
                else:
                    label_ids.append(-100)
                previous_word_idx = word_idx
            labels.append(label_ids)
        tokenized_inputs["labels"] = labels
    return tokenized_inputs

train_dataset = train_dataset.map(tokenize_and_align_labels, batched=True, )
test_dataset = test_dataset.map(tokenize_and_align_labels, batched=True, )
```

Handwritten note: 256 이외의 것들은 잘라주겠다

256 이외의 것들은 잘라주겠다

How to train NER using transformers

- Preprocessing the dataset – set tag ids

tokenized 하고
label 변경 해주는 작업

Raw text	the		adenomatous			polyposis			coli	tumour	
label ids	0		1			2			2	2	
Tokenized	the	aden	##oma	##tou	##s	poly	##po	##sis	coli	tu	##mour
label ids	0	1	-100	-100	-100	2	-100	-100	2	2	-100

go having

going → ?

Tag ID	Tag
1	B
2	I
-100	<sub_token>

How to train NER using transformers

- Set metric function

성능 평가

transformer에 같이 넣어서

```
def compute_metrics(p):  
    metric = evaluate.load("segeval")  
  
    predictions, labels = p  
    predictions = np.argmax(predictions, axis=2)  
  
    # Remove ignored index (special tokens)  
    true_predictions = [  
        [label_list[p] for (p, l) in zip(prediction, label) if l != -100]  
        for prediction, label in zip(predictions, labels)  
    ]  
    true_labels = [  
        [label_list[l] for (p, l) in zip(prediction, label) if l != -100]  
        for prediction, label in zip(predictions, labels)  
    ]  
  
    results = metric.compute(predictions=true_predictions, references=true_labels)  
    return {  
        "precision": results["overall_precision"],  
        "recall": results["overall_recall"],  
        "f1": results["overall_f1"],  
        "accuracy": results["overall_accuracy"],  
    }
```

성능 평가에 필요한 애들만 모음

How to train NER using transformers


- Set metric function

$$precision = \frac{TP}{TP + FP}$$


$$recall = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2 \times precision \times recall}{precision + recall}$$


$$Accuracy = \frac{TP + TN}{TP + FN + TN + FP}$$



How many retrieved items are relevant?



How many relevant items are retrieved?



How to train NER using transformers


- Set metric function

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2 \times precision \times recall}{precision + recall}$$

$$Accuracy = \frac{TP + TN}{TP + FN + TN + FP}$$



Identification of ACP2 , a homologue of the adenomatous polyposis coli tumour suppressor .

O O O O O O O O B I I I O O

positive

How to train NER using transformers

- Train NER model

Dataset 학습시 배치 사이즈로 묶어 학습함

```
data_collator = DataCollatorForTokenClassification(tokenizer, )
trainer = Trainer(
    model=model, Bert
    args=training_args,
    train_dataset=train_dataset,
    data_collator=data_collator,
    compute_metrics=compute_metrics
)

train_result = trainer.train()
metrics = train_result.metrics
print(metrics)
```

How to train NER using transformers

- Train NER model

Step Training Loss

500	0.174000
1000	0.108500
1500	0.112300
2000	0.098300
2500	0.078000
3000	0.078300
3500	0.055800
4000	0.058400
4500	0.053600
5000	0.057000

Training completed. Do not forget to share your model on huggingface.co/models =)

Saving model checkpoint to ./results

Configuration saved in ./results/config.json


Model weights saved in ./results/pytorch_model.bin

{'train_runtime': 522.341, 'train_samples_per_second': 10.401, 'train_steps_per_second': 10.401, 'total_flos': 709819057046016.0, 'train_loss': 0.08363902026856997, 'epoch': 1.0}

How to train NER using transformers

```
%load_ext tensorboard  
%tensorboard --logdir="./results"
```

train/loss



How to train NER using transformers

- Test NER model

```
predictions, labels, metrics = trainer.predict(test_dataset, metric_key_prefix="predict")
print(metrics)

predictions = np.argmax(predictions, axis=2) 가장 높은 확률을 가지고 옴

true_predictions = [
    [label_list[p] for (p, l) in zip(prediction, label) if l != -100]
    for prediction, label in zip(predictions, labels)
]

for token, prediction in zip(test_dataset["tokens"][0], true_predictions[0]):
    print(f"{token} \t {prediction}")
```

How to train NER using transformers

- Test NER model

```
***** Running Prediction *****
```

```
Num examples = 941
```

```
Batch size = 11
```

```
*****Metrics*****
```

```
predict_loss: 0.07623732835054398
```

```
predict_precision: 0.8138613861386138
```

```
predict_recall: 0.85625
```

```
predict_f1: 0.8345177664974618
```

```
predict_accuracy: 0.9808956198718211
```

```
predict_runtime: 21.2564
```

```
predict_samples_per_second: 44.269
```

```
predict_steps_per_second: 4.046
```

```
****Result****
```

```
Clustering      0
```

```
of              0
```

```
missense       0
```

```
mutations      0
```

```
in             0
```

```
the            0
```

```
ataxia B-Disease
```

```
-        I-Disease
```

```
telangiectasia I-Disease
```

```
gene          0
```

```
in            0
```

```
a             0
```

```
sporadic      0
```

```
T            B-Disease
```

```
-            I-Disease
```

```
cell          I-Disease
```

```
leukaemia     I-Disease
```

```
.             0
```

Ataxia-telangiectasia is a rare inherited disorder that affects the nervous system, immune system, and other body systems.

T-cell leukemia is an uncommon type of blood cell cancer that affects your white blood cells.